

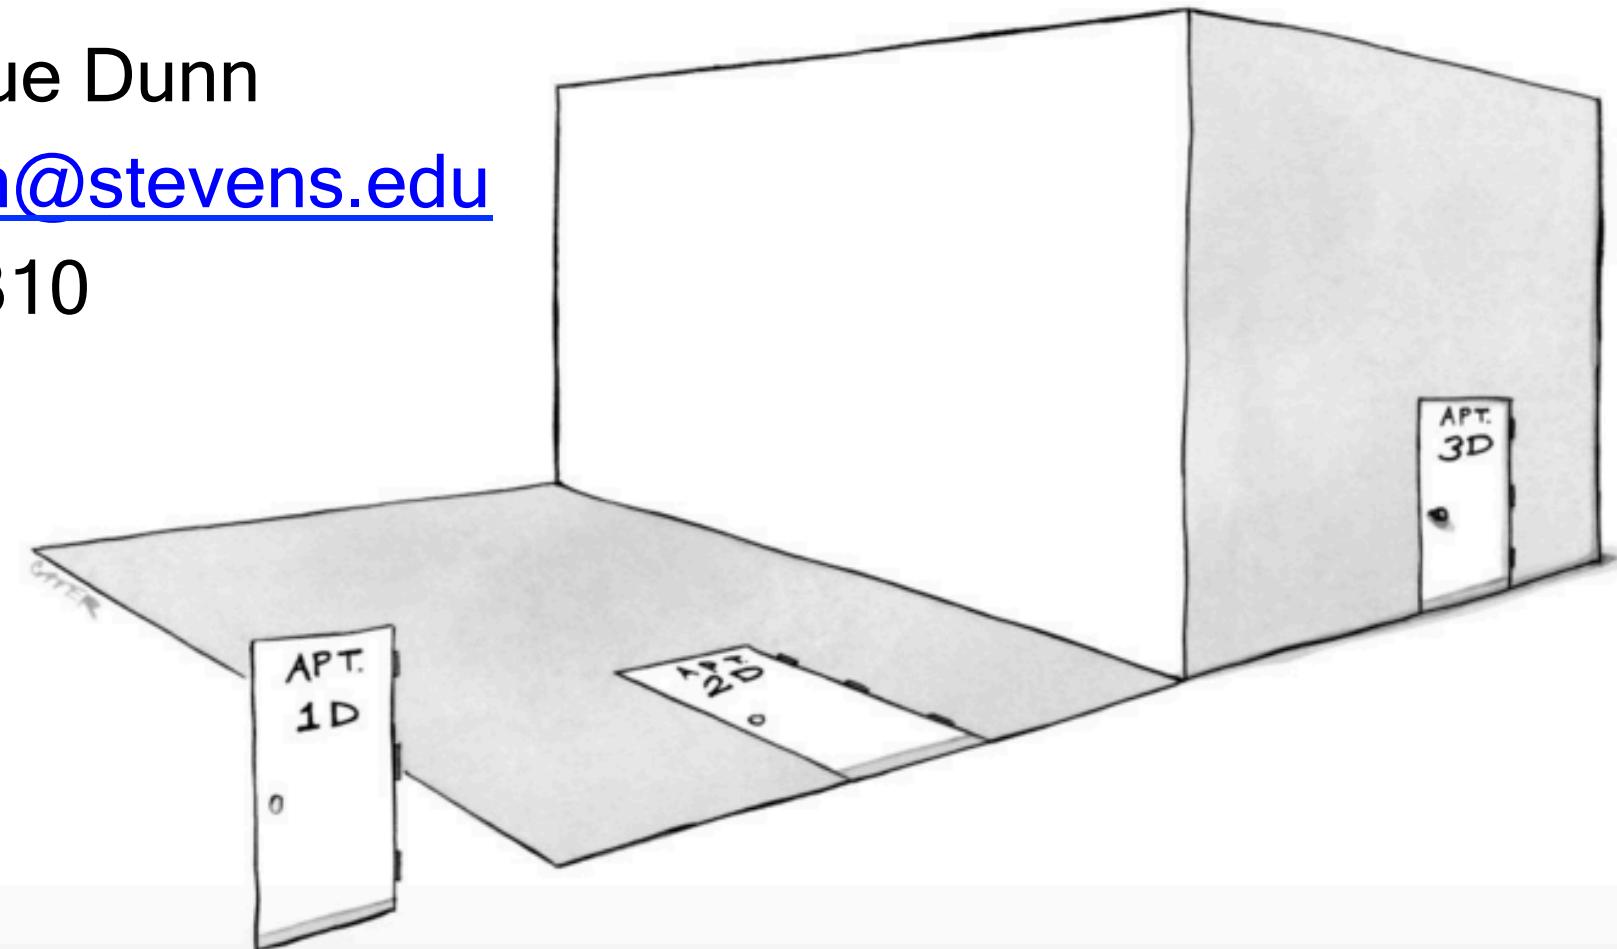
# CS 532: 3D Computer Vision

## Lecture 3

Enrique Dunn

[edunn@stevens.edu](mailto:edunn@stevens.edu)

Lieb 310



# Course TA

Andy Wiggins

<awiggins@stevens.edu>

Office hours:

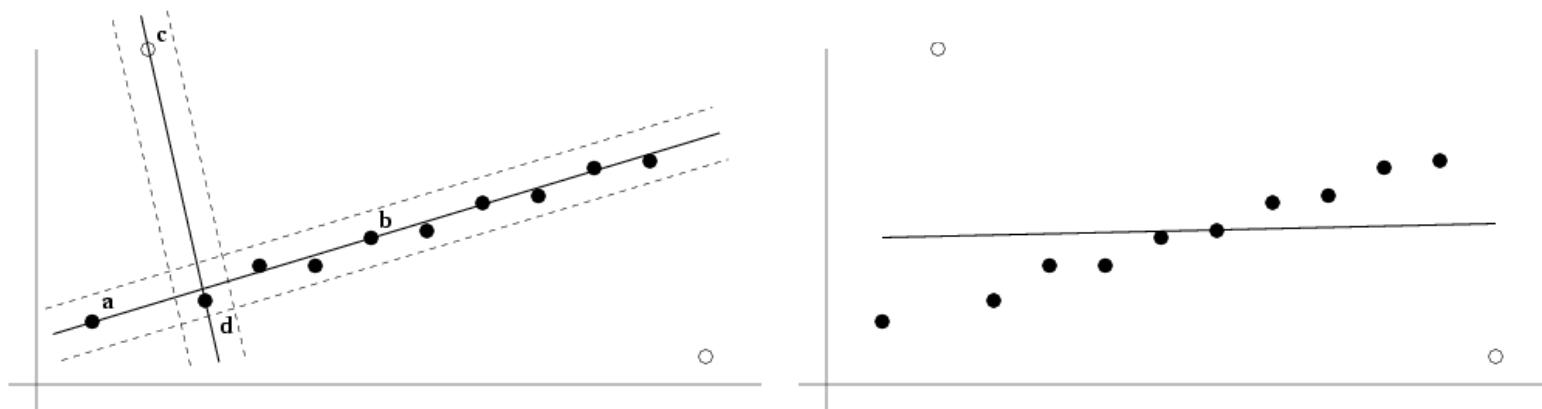
Lieb lounge on Wednesdays & Thursdays 2pm-4pm

# RANSAC

Slides by R. Hartley, A. Zisserman  
and M. Pollefeys

# Robust Estimation

- What if set of matches contains gross outliers?



# RANSAC

## Objective

Robust fit of model to data set  $S$  which contains outliers

## Algorithm

- (i) Randomly select a sample of  $s$  data points from  $S$  and instantiate the model from this subset.
- (ii) Determine the set of data points  $S_i$  which are within a **distance threshold**  $t$  of the model. The set  $S_i$  is the consensus set of samples and defines the inliers of  $S$ .
- (iii) If the subset of  $S_i$  is greater than some threshold  $T$ , re-estimate the model using all the points in  $S_i$  and terminate
- (iv) If the size of  $S_i$  is less than  $T$ , select a new subset and repeat the above.
- (v) After  $N$  trials the largest consensus set  $S_i$  is selected, and the model is re-estimated using all the points in the subset  $S_i$

# How Many Samples?

Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers. e.g.  $p=0.99$

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

Sampling Inlier Data point

$$1 - e$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

Sampling All-Inlier Set

$$(1 - e)^s$$

Sampling Contaminated Set

$$\frac{1}{1 - (1 - e)^s}$$

s	proportion of outliers $e$							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

# Acceptable Consensus Set

- Typically, terminate when inlier ratio reaches expected ratio of inliers

$$T = (1 - e)n$$

# Adaptively Determining the Number of Samples

$e$  is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield  $e=0.2$

- $N=\infty$ ,  $\text{sample\_count}=0$
  - While  $N > \text{sample\_count}$  repeat
    - Choose a sample and count the number of inliers
    - Set  $e = 1 - (\text{number of inliers}) / (\text{total number of points})$
    - Recompute  $N$  from  $e$
    - Increment the  $\text{sample\_count}$  by 1
  - Terminate
- $$(N = \log(1 - p) / \log(1 - (1 - e)^s))$$

# Other robust algorithms

- RANSAC maximizes number of inliers
- LMedS minimizes median error
- Not recommended: case deletion, iterative least-squares, etc.

# Automatic Computation of H

## Objective

Compute homography between two images

## Algorithm

- (i) **Interest points:** Compute interest points in each image
- (ii) **Putative correspondences:** Compute a set of interest point matches based on some similarity measure
- (iii) **RANSAC robust estimation:** Repeat for  $N$  samples
  - (a) Select 4 correspondences and compute H
  - (b) Calculate the distance  $d_{\perp}$  for each putative match
  - (c) Compute the number of inliers consistent with H ( $d_{\perp} < t$ )Choose H with most inliers
- (iv) **Optimal estimation:** re-estimate H from all inliers by minimizing ML cost function with Levenberg-Marquardt
- (v) **Guided matching:** Determine more matches using prediction by computed H

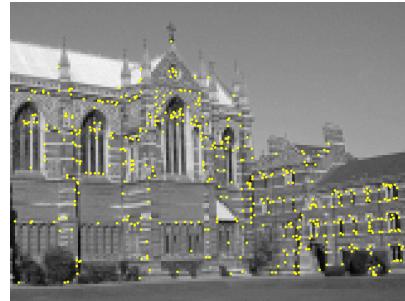
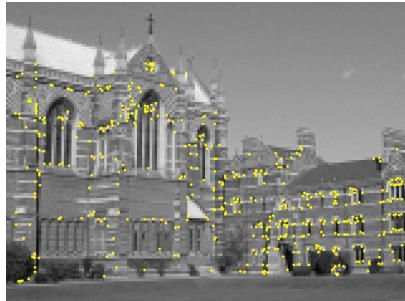
Optionally iterate last two steps until convergence

# Determine Putative Correspondences

- Compare interest points
  - Similarity measure:
    - SAD, SSD, ZNCC in small neighborhood
- If motion is limited, only consider interest points with similar coordinates

# Example: robust computation

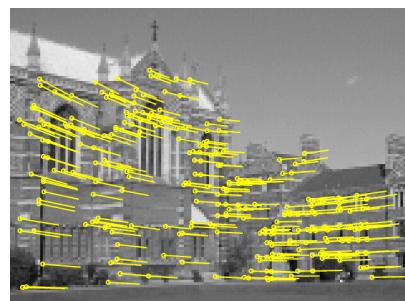
#in	1-e	adapt. $N$
6	2%	20M
10	3%	2.5M
44	16%	6,922
58	21%	2,291
73	26%	911
<u>151</u>	<u>56%</u>	<u>43</u>



Interest points  
(500/image)  
(640x480)



Putative correspondences (268)  
(Best match, SSD < 20)  
Outliers (117)  
( $t=1.25$  pixel; 43 iterations)



Inliers (151)

Final inliers (262)

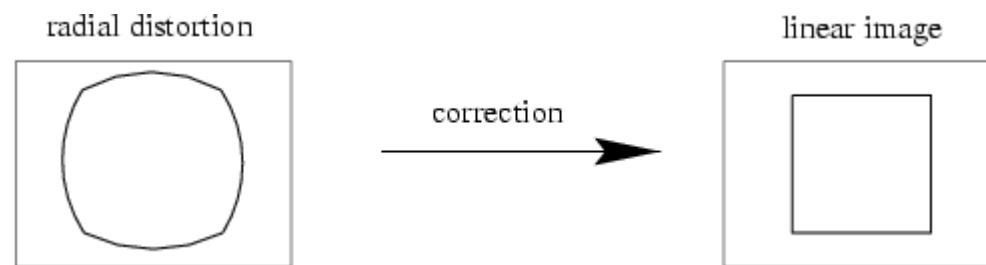
# Radial Distortion and Undistortion

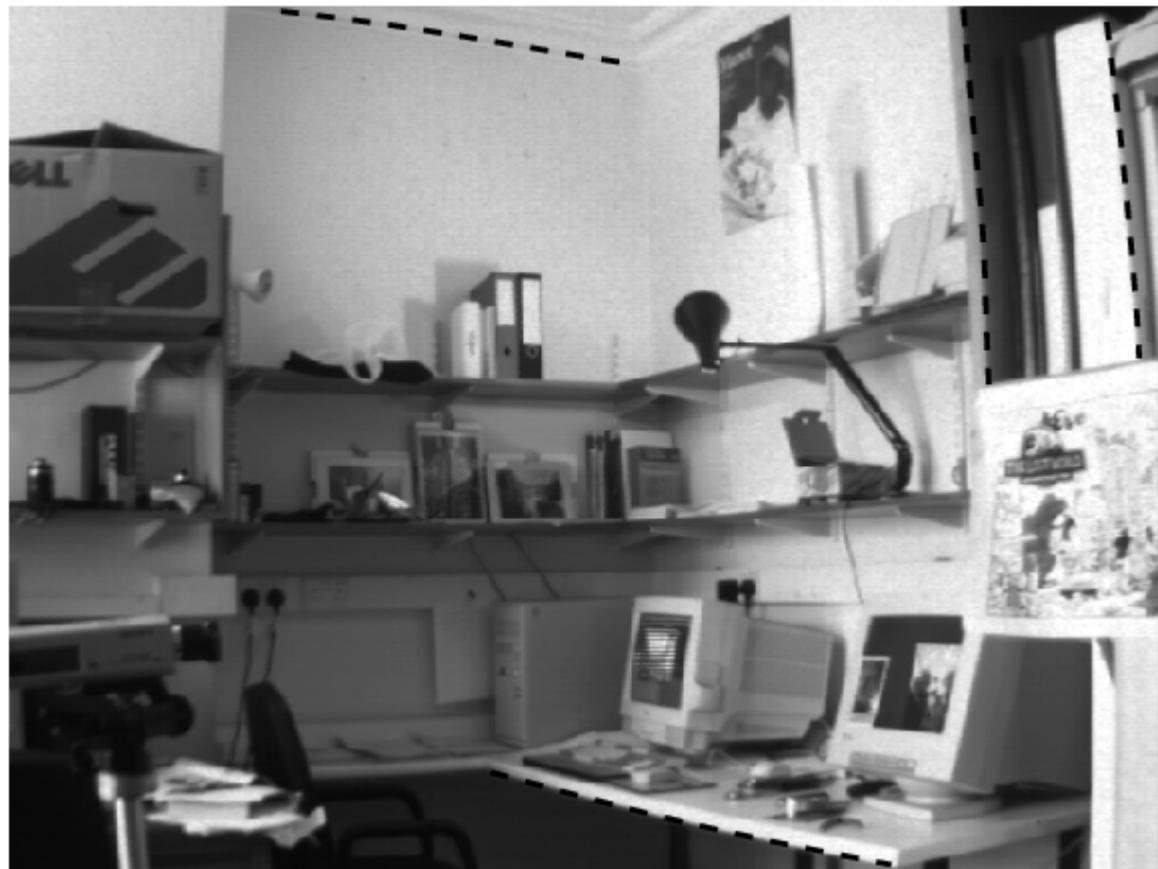
Slides by R. Hartley, A. Zisserman  
and M. Pollefeys

# Radial Distortion



short and long focal length







# Typical Undistortion Model

Correction of distortion

$$\hat{x} = x_c + L(r)(x - x_c) \quad \hat{y} = y_c + L(r)(y - y_c)$$

Choice of the distortion function and center

$$x = x_o + (x_o - c_x)(K_1 r^2 + K_2 r^4 + \dots)$$
$$y = y_o + (y_o - c_y)(K_1 r^2 + K_2 r^4 + \dots)$$

$$r = (x_o - c_x)^2 + (y_o - c_y)^2 .$$

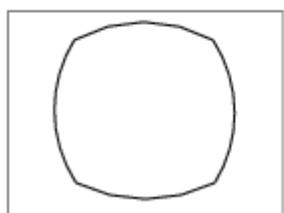
Computing the parameters of the distortion function

- (i) Minimize with additional unknowns
- (ii) Straighten lines
- (iii) ...

# Why Undistort?

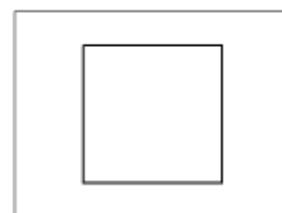


radial distortion



→  
correction

linear image



$$(\tilde{x}, \tilde{y}, 1)^\top = [\mathbf{I} \mid \mathbf{0}] \mathbf{X}_{\text{cam}}$$

$$\begin{pmatrix} x_d \\ y_d \end{pmatrix} = L(\tilde{r}) \begin{pmatrix} \tilde{x} \\ \tilde{y} \end{pmatrix}$$

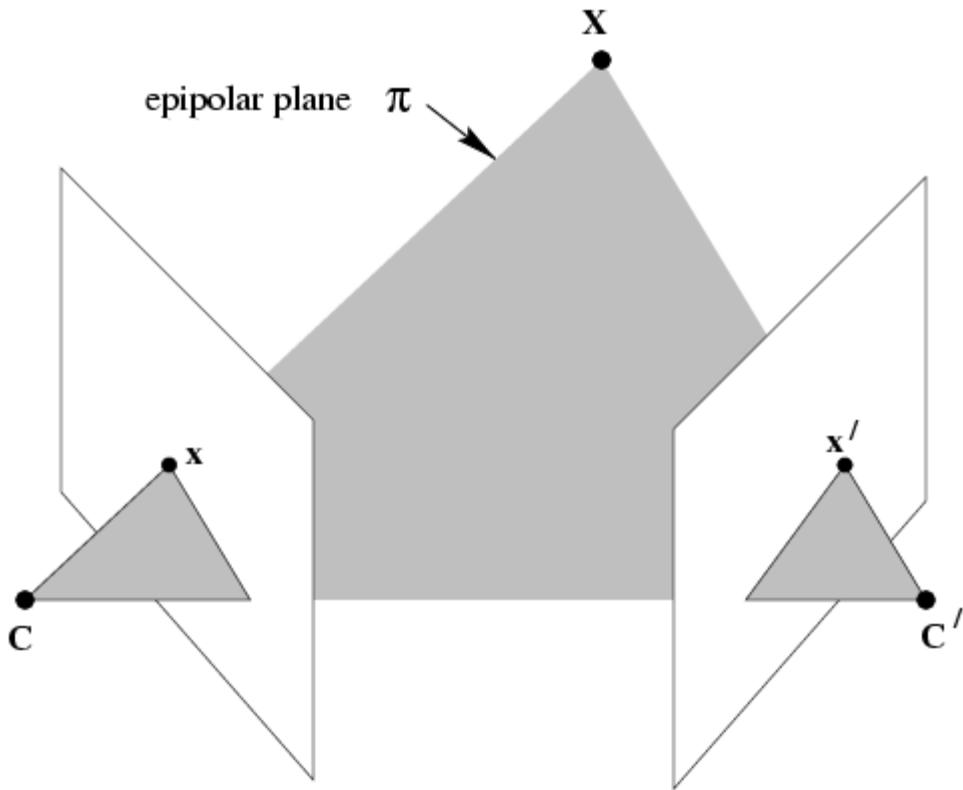
# Two-View Geometry

Slides by R. Hartley, A. Zisserman and M. Pollefeys

# Three questions:

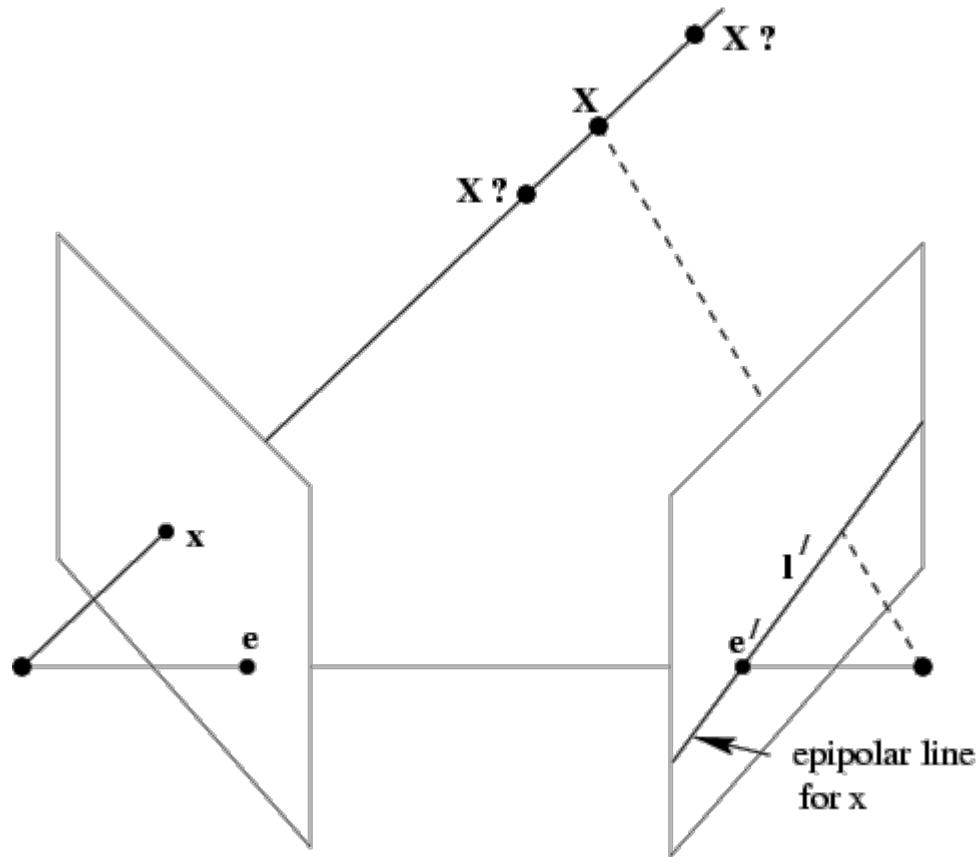
- (i) **Correspondence geometry:** Given an image point  $x$  in the first image, how does this constrain the position of the corresponding point  $x'$  in the second image?
- (ii) **Camera geometry (motion):** Given a set of corresponding image points  $\{x_i \leftrightarrow x'_i\}$ ,  $i=1,\dots,n$ , what are the cameras  $P$  and  $P'$  for the two views?
- (iii) **Scene geometry (structure):** Given corresponding image points  $x_i \leftrightarrow x'_i$  and cameras  $P, P'$ , what is the position of (their pre-image)  $X$  in space?

# The Epipolar Geometry



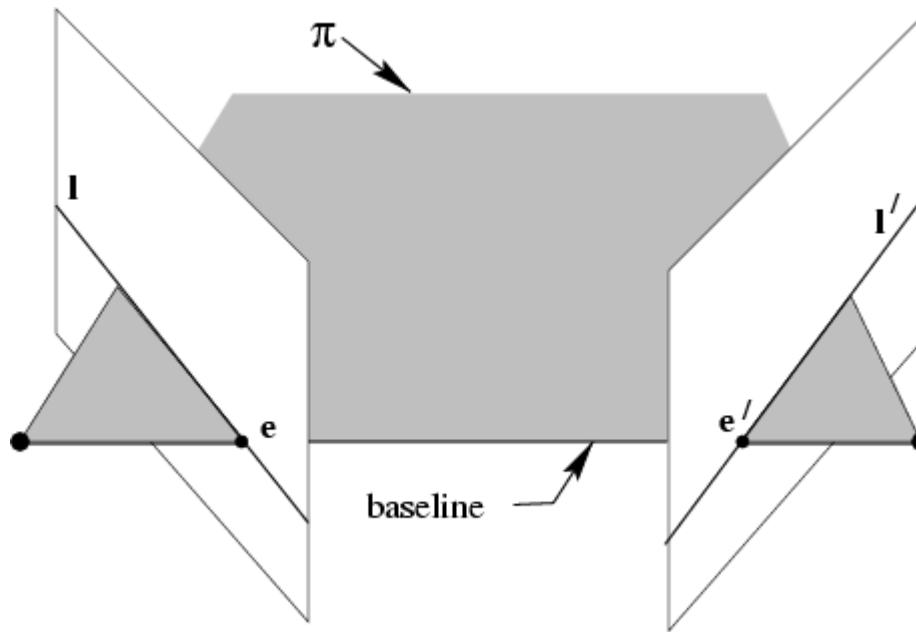
$C$ ,  $C'$ ,  $x$ ,  $x'$  and  $X$  are coplanar

# The Epipolar Geometry



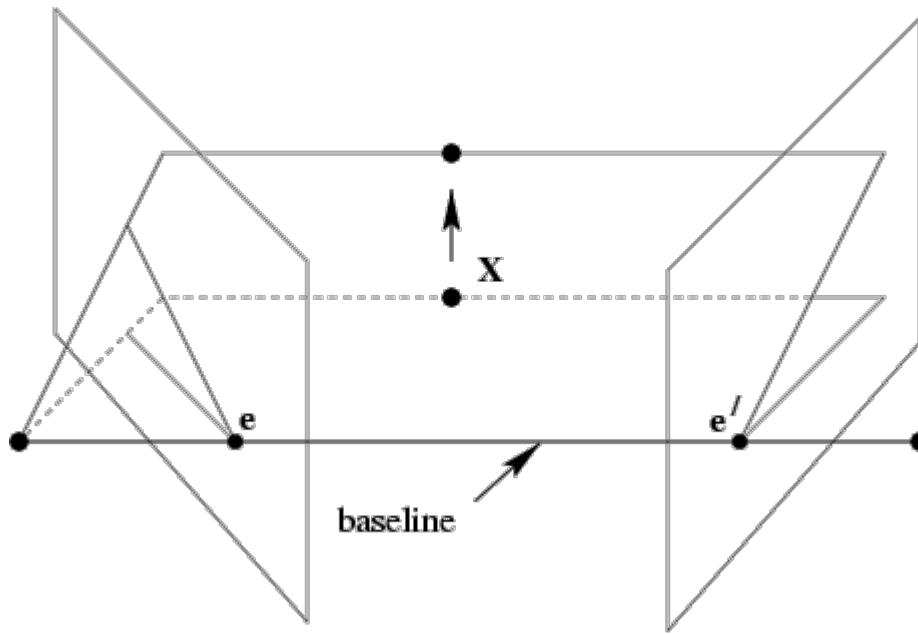
What if only  $C, C', x$  are known?

# The Epipolar Geometry



All points on  $\pi$  project on  $l$  and  $l'$

# The Epipolar Geometry



Family of planes  $\pi$  and lines  $l$  and  $l'$   
Intersection in  $e$  and  $e'$

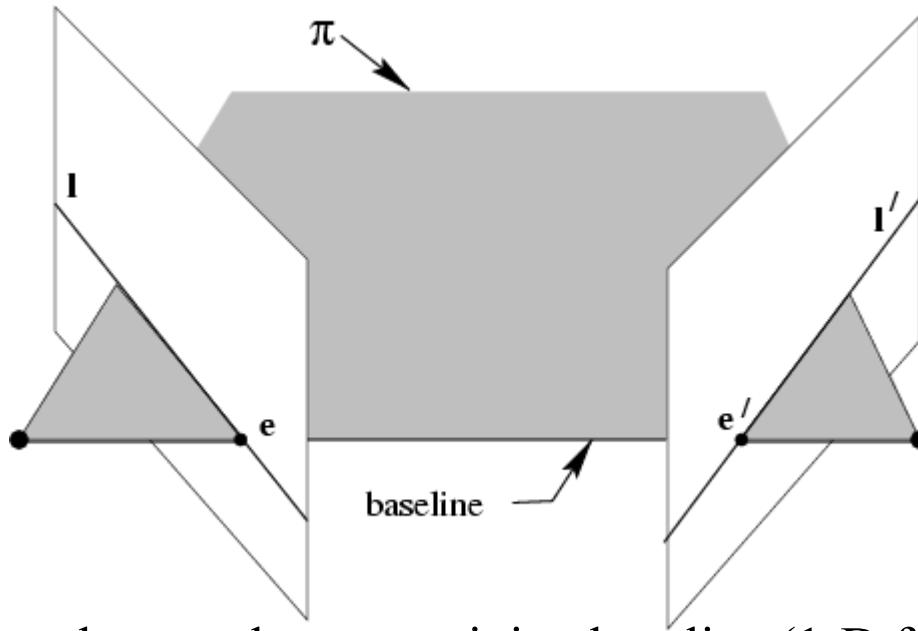
# The Epipolar Geometry

epipoles  $e, e'$

= intersection of baseline with image plane

= projection of projection center in other image

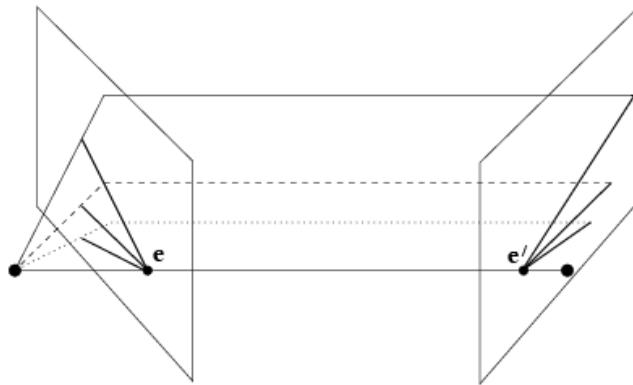
= vanishing point of camera motion direction



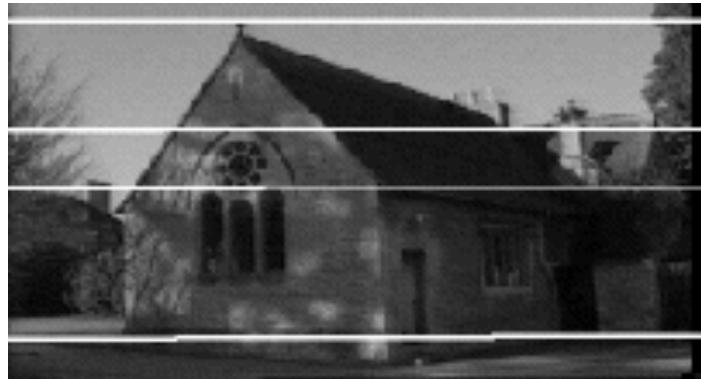
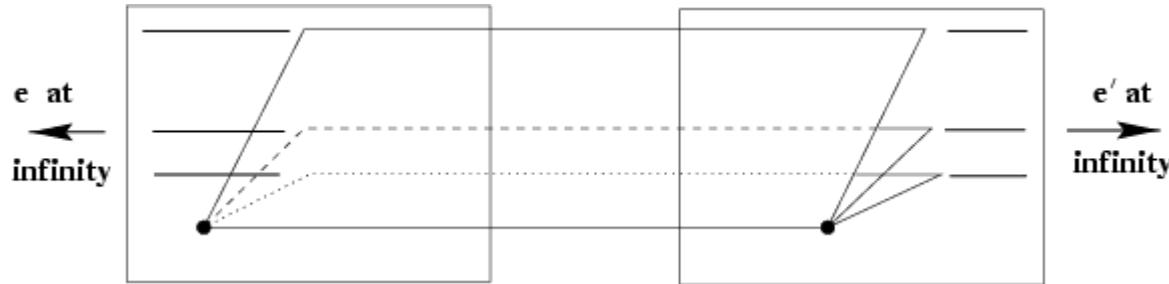
an epipolar plane = plane containing baseline (1-D family)

an epipolar line = intersection of epipolar plane with image  
(always come in corresponding pairs)

# Example: Converging Cameras

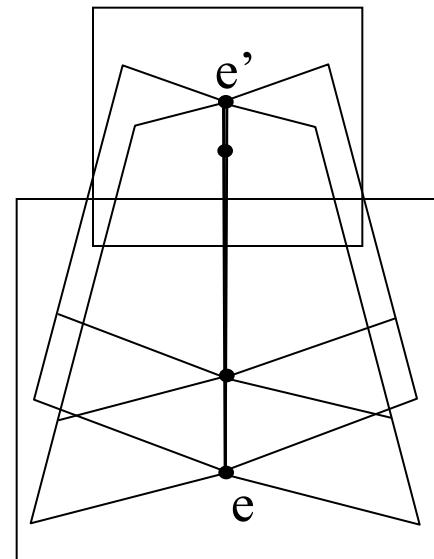
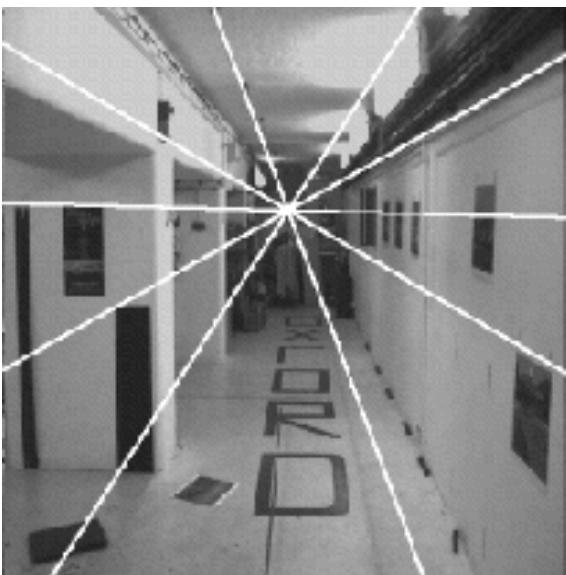
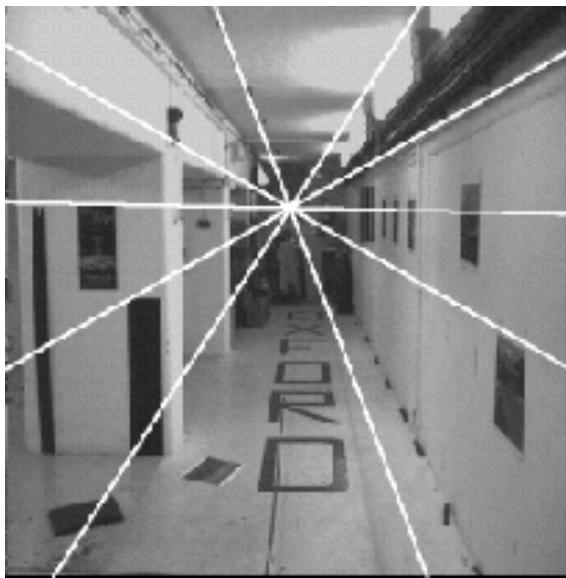


# Example: Motion Parallel to Image Plane



(simple for stereo → rectification)

# Example: Forward Motion



# The Fundamental Matrix $F$

algebraic representation of epipolar geometry

$$x \mapsto l'$$

we will see that mapping is a (singular) correlation  
(i.e. projective mapping from points to lines)  
represented by the fundamental matrix  $F$

# The Fundamental Matrix F

correspondence condition

The fundamental matrix satisfies the condition that for any pair of corresponding points  $x \leftrightarrow x'$  in the two images

$$x'^T F x = 0 \quad (x'^T l' = 0)$$

# The Fundamental Matrix F

F is the unique  $3 \times 3$  rank 2 matrix that satisfies  $x'^T F x = 0$  for all  $x \leftrightarrow x'$

- (i) **Transpose:** if F is fundamental matrix for  $(P, P')$ , then  $F^T$  is fundamental matrix for  $(P', P)$
- (ii) **Epipolar lines:**  $l' = Fx$  &  $l = F^T x'$
- (iii) **Epipoles:** on all epipolar lines, thus  $e'^T F x = 0$ ,  $\forall x \Rightarrow e'^T F = 0$ , similarly  $F e = 0$
- (iv) F has 7 d.o.f. , i.e.  $3 \times 3 - 1$  (homogeneous) - 1 (rank 2)
- (v) F is a correlation, projective mapping from a point x to a line  $l' = Fx$  (not a proper correlation, i.e. not invertible)

## Two View Geometry Computation: Linear Algorithm

For every match  $(m, m')$ :  $x'^T F x = 0$

$$x'xf_{11} + x'yf_{12} + x'f_{13} + y'xf_{21} + y'yf_{22} + y'f_{23} + xf_{31} + yf_{32} + f_{33} = 0$$

separate known from unknown

$$\begin{bmatrix} x'_1 x_1 & x'_1 y_1 & x'_1 & y'_1 x_1 & y'_1 y_1 & y'_1 & x_1 & y_1 & 1 \\ \vdots & \vdots \\ x'_n x_n & x'_n y_n & x'_n & y'_n x_n & y'_n y_n & y'_n & x_n & y_n & 1 \end{bmatrix} f = 0$$

$$Af = 0$$

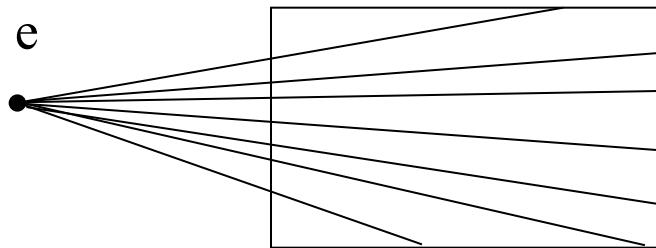
# Benefits from having $F$

- Given a pixel in one image, the corresponding pixel has to lie on epipolar line
- Search space reduced from 2-D to 1-D

# Image Pair Rectification

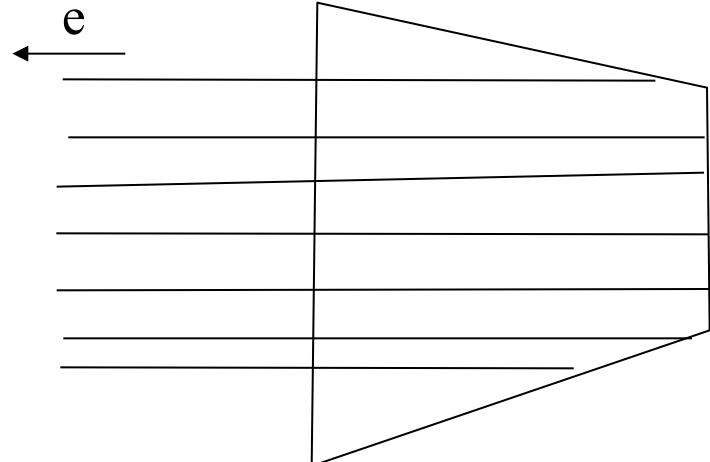
simplify stereo matching  
by warping the images

Apply projective transformation so that epipolar lines  
correspond to horizontal scanlines



$$\begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = H\mathbf{e}$$

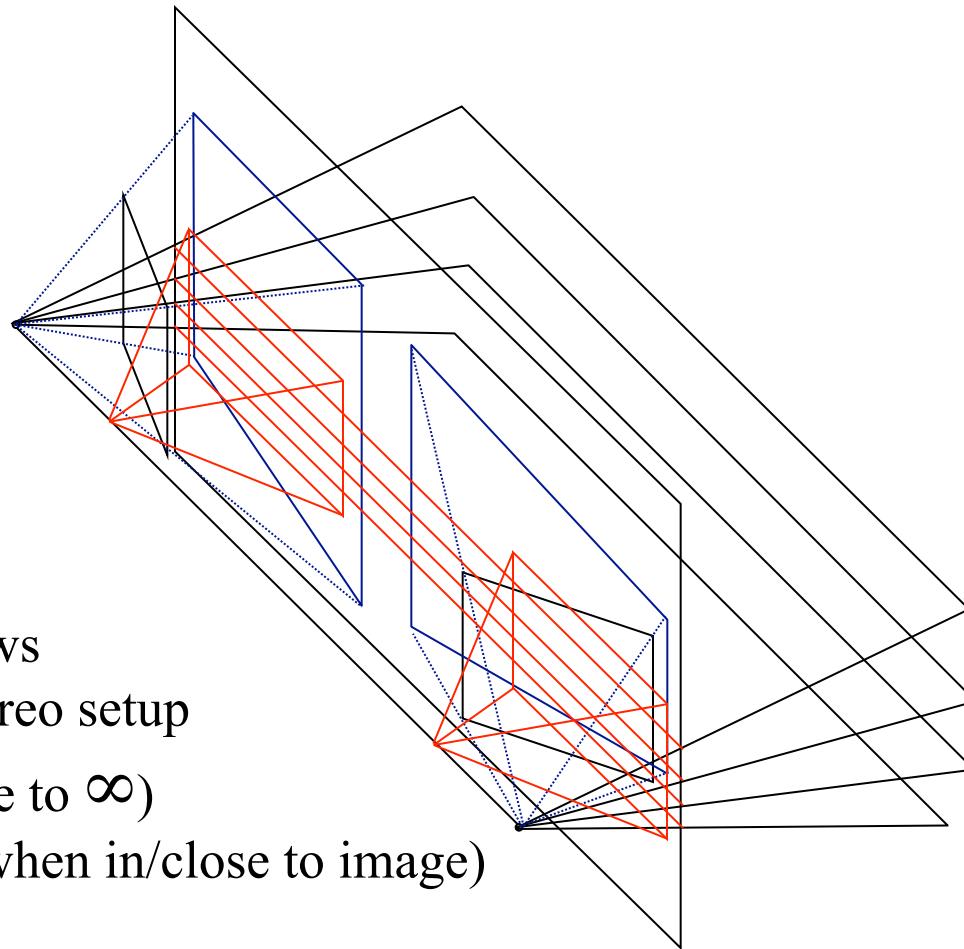
map epipole  $e$  to  $(1,0,0)$   
try to minimize image distortion



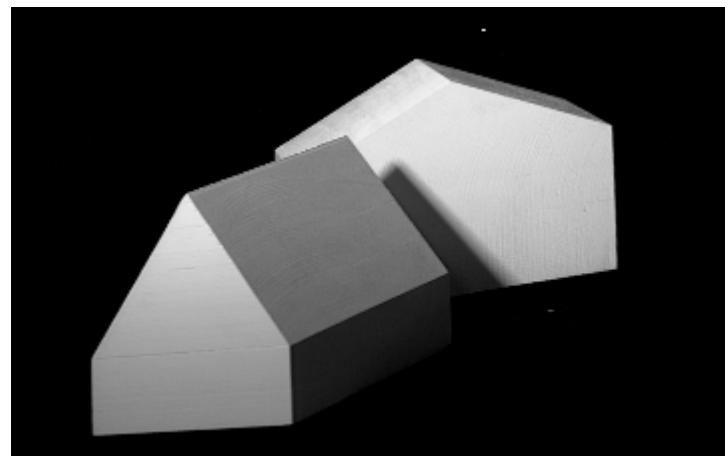
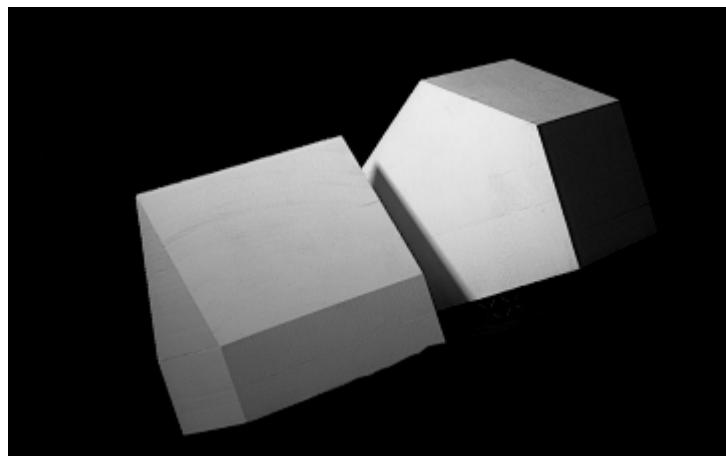
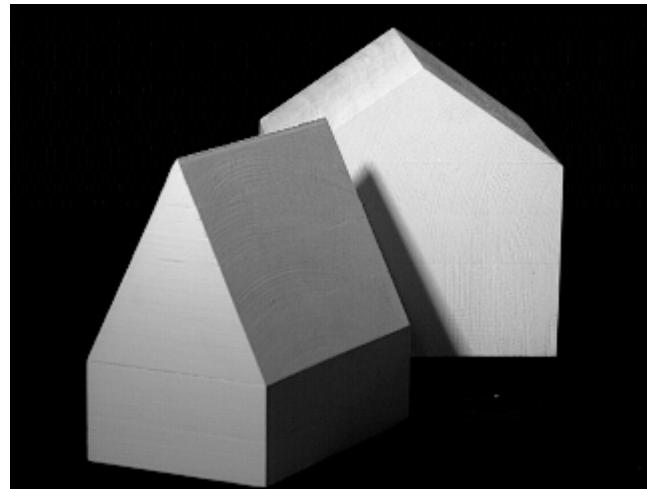
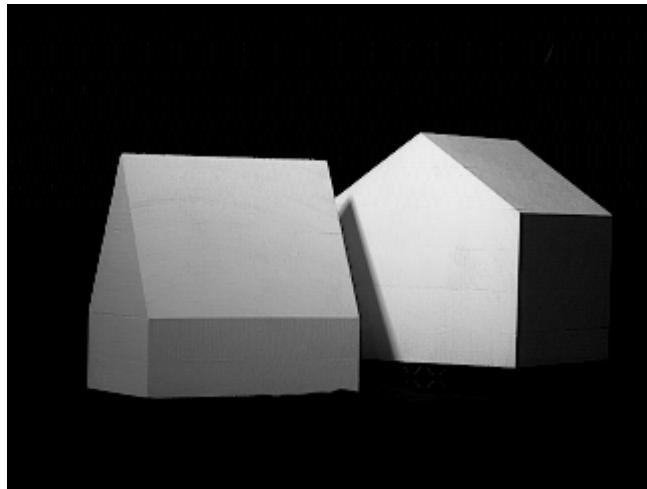
problem when epipole in (or close to) the image

# Planar Rectification

(standard approach)



Bring two views  
to standard stereo setup  
(moves epipole to  $\infty$ )  
(not possible when in/close to image)



# The Essential Matrix

~fundamental matrix for calibrated cameras (remove K)

$$E = [t]_x R = R[R^T t]_x \quad [\mathbf{a}]_x \stackrel{\text{def}}{=} \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}.$$

$$\hat{x}'^T E \hat{x} = 0 \quad (\hat{x} = K^{-1}x; \hat{x}' = K^{-1}x')$$

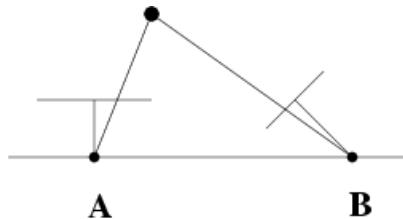
$$E = K'^T F K$$

5 d.o.f. (3 for R; 2 for t up to scale)

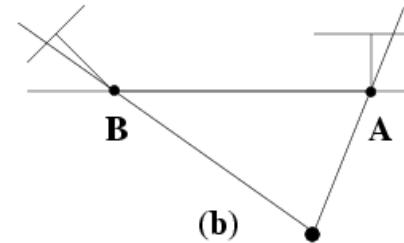
E is an essential matrix if and only if two singular values are equal (and the third=0)

$$E = U \text{diag}(1, 1, 0) V^T$$

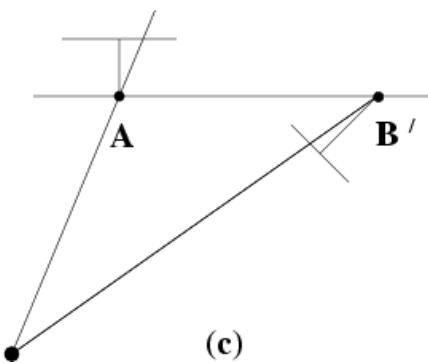
# Four Possible Solutions from E



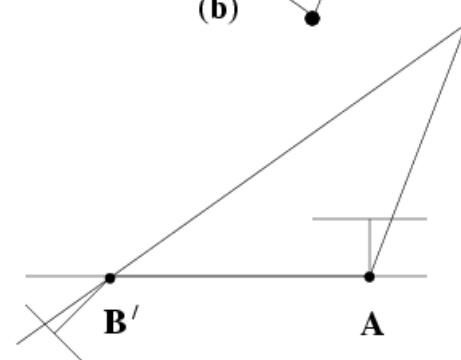
(a)



(b)



(c)



(d)

Given E and setting the first camera matrix  $P = [I \mid 0]$ , there are four possible solutions for  $P'$  (only one solution, however, where a reconstructed point is in front of both cameras)

# Fundamental Matrix Estimation

# Epipolar Geometry: Basic Equation

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

$$x'xf_{11} + x'yf_{12} + x'f_{13} + y'xf_{21} + y'yf_{22} + y'f_{23} + xf_{31} + yf_{32} + f_{33} = 0$$

## separate known from unknown

$$\begin{bmatrix} x'_1 & x_1 & x'_1 & y'_1 & x_1 & y'_1 & y_1 & x_1 & y_1 & 1 \\ \vdots & \vdots \\ x'_n & x_n & x'_n & y'_n & x_n & y'_n & y_n & x_n & y_n & 1 \end{bmatrix} f = 0$$

$$Af = 0$$

# The Singularity Constraint

$$e'^T F = 0 \quad Fe = 0 \quad \det F = 0 \quad \text{rank } F = 2$$

SVD from linearly computed F matrix (rank 3)

$$F = U \begin{bmatrix} \sigma_1 & & \\ & \sigma_2 & \\ & & \sigma_3 \end{bmatrix} V^T = U_1 \sigma_1 V_1^T + U_2 \sigma_2 V_2^T + U_3 \sigma_3 V_3^T$$

Compute closest rank-2 approximation  $\min \|F - F'\|_F$

$$F' = U \begin{bmatrix} \sigma_1 & & \\ & \sigma_2 & \\ & & 0 \end{bmatrix} V^T = U_1 \sigma_1 V_1^T + U_2 \sigma_2 V_2^T$$

# The Singularity Constraint



# The NOT Normalized 8-point Algorithm

$$\begin{bmatrix} x_1x_1' & y_1x_1' & x_1' & x_1y_1' & y_1y_1' & y_1' & x_1 & y_1 & 1 \\ x_2x_2' & y_2x_2' & x_2' & x_2y_2' & y_2y_2' & y_2' & x_2 & y_2 & 1 \\ \vdots & \vdots \\ x_nx_n' & y_nx_n' & x_n' & x_ny_n' & y_ny_n' & y_n' & x_n & y_n & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix} = 0$$

~10000

~10000

~100

~10000

~10000

~100

~100

~100

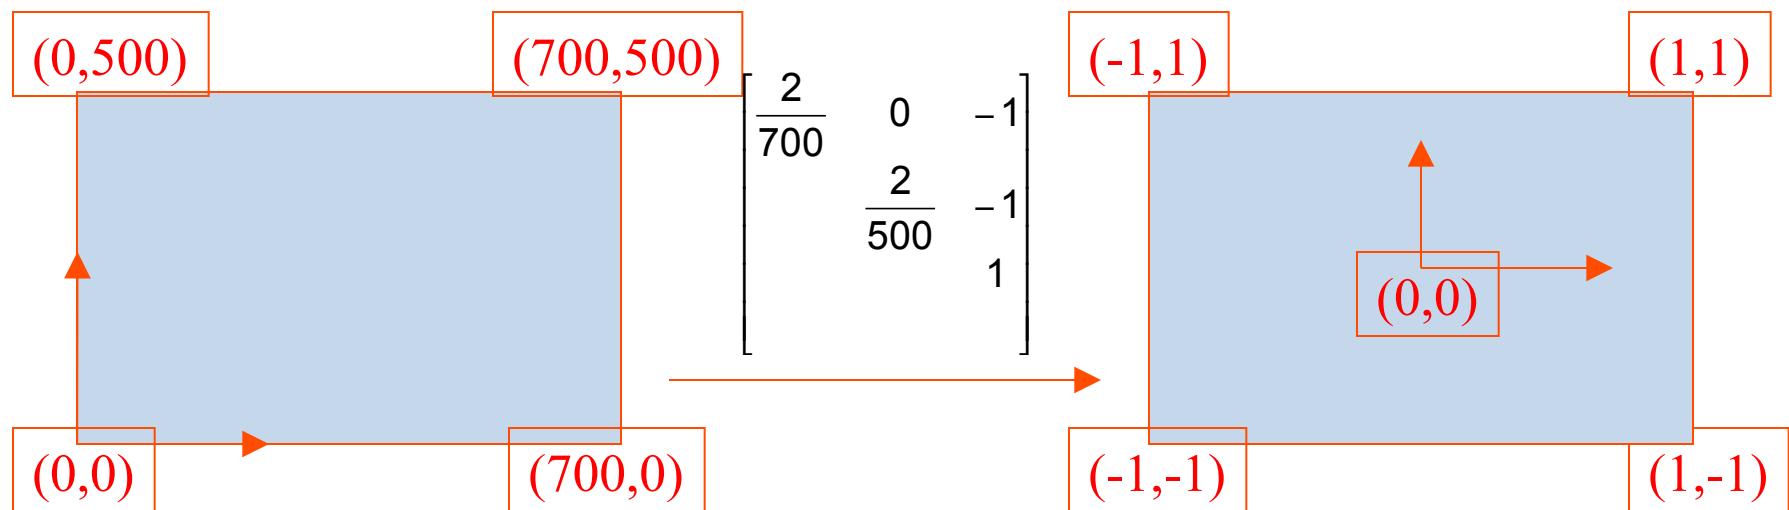
1



Orders of magnitude difference  
between column of data matrix  
→ least-squares yields poor results

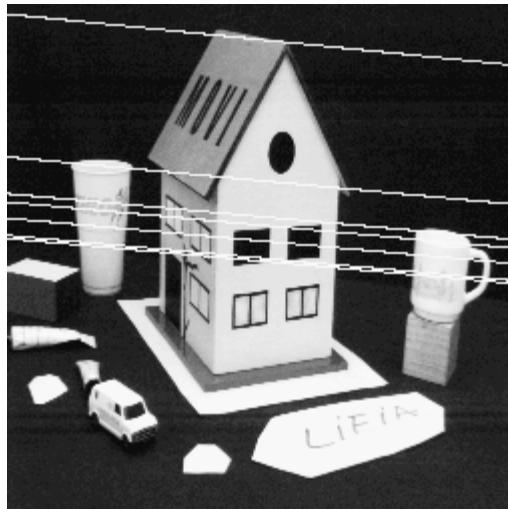
# The Normalized 8-point Algorithm

Transform image to  $[-1,1] \times [-1,1]$

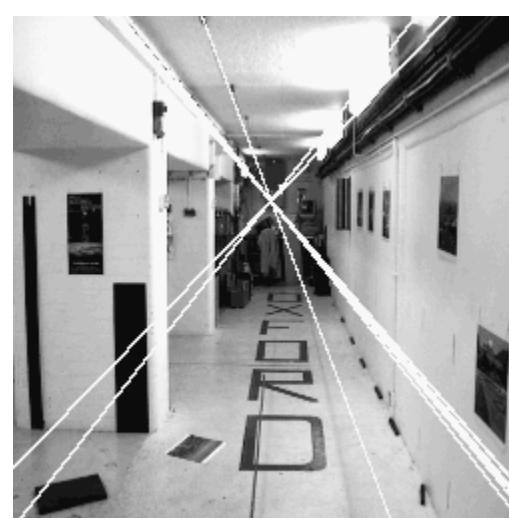
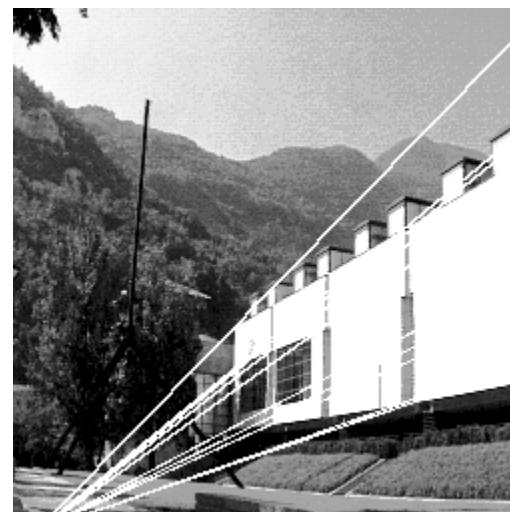
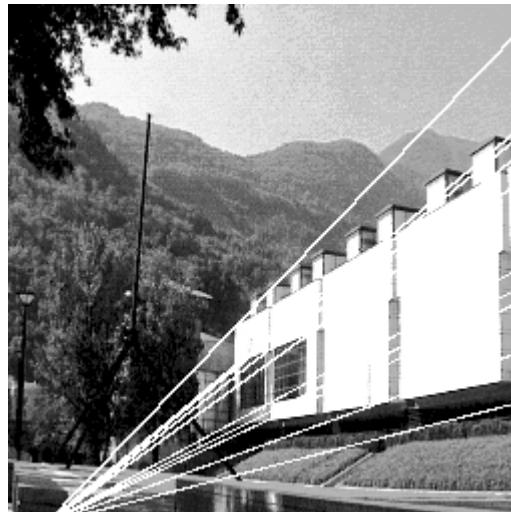


normalized least squares yields good results

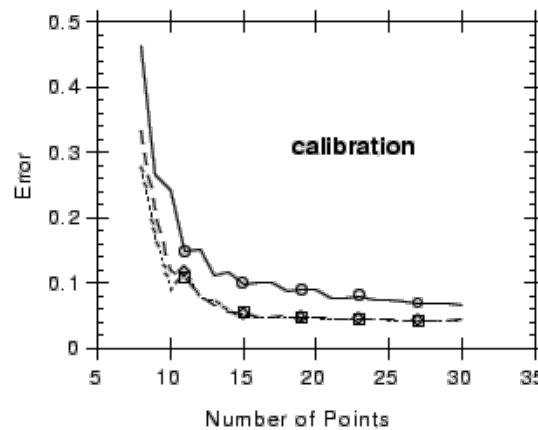
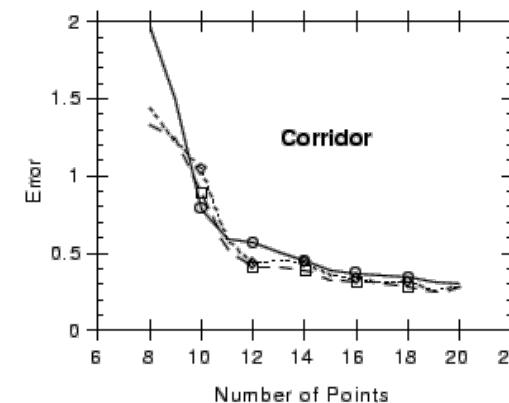
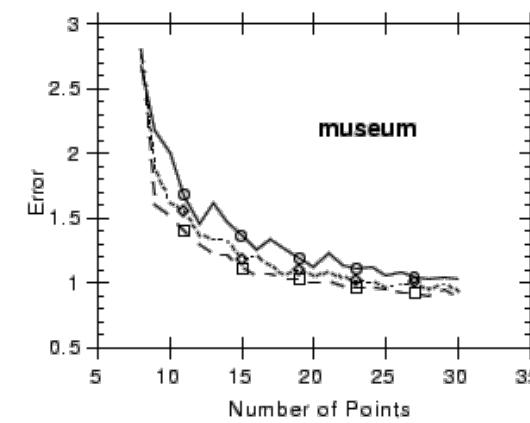
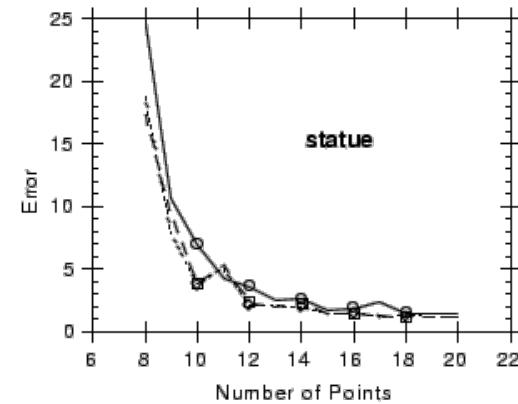
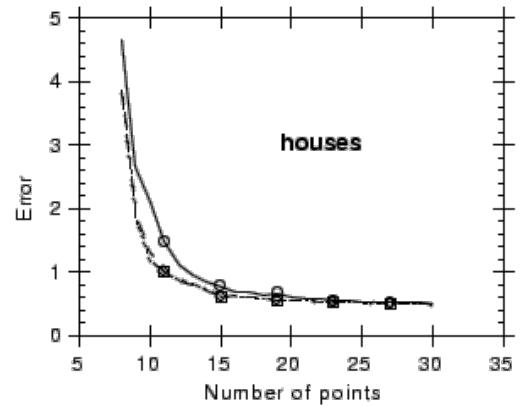
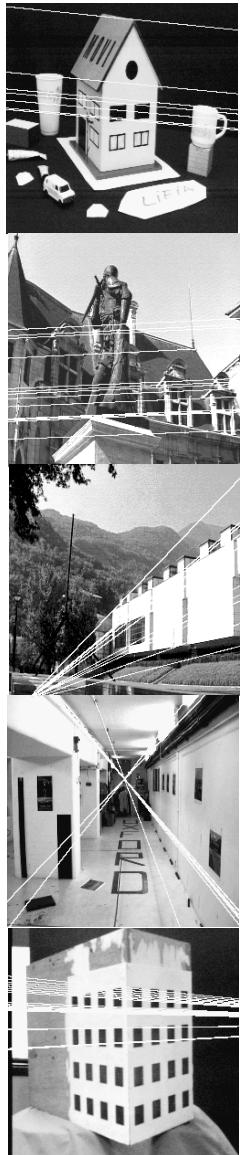
# Some Experiments



# Some Experiments



# Some Experiments



Residual error:

$$\sum_i d(\mathbf{x}'_i, F\mathbf{x}_i)^2 + d(\mathbf{x}_i, F^T \mathbf{x}'_i)^2$$

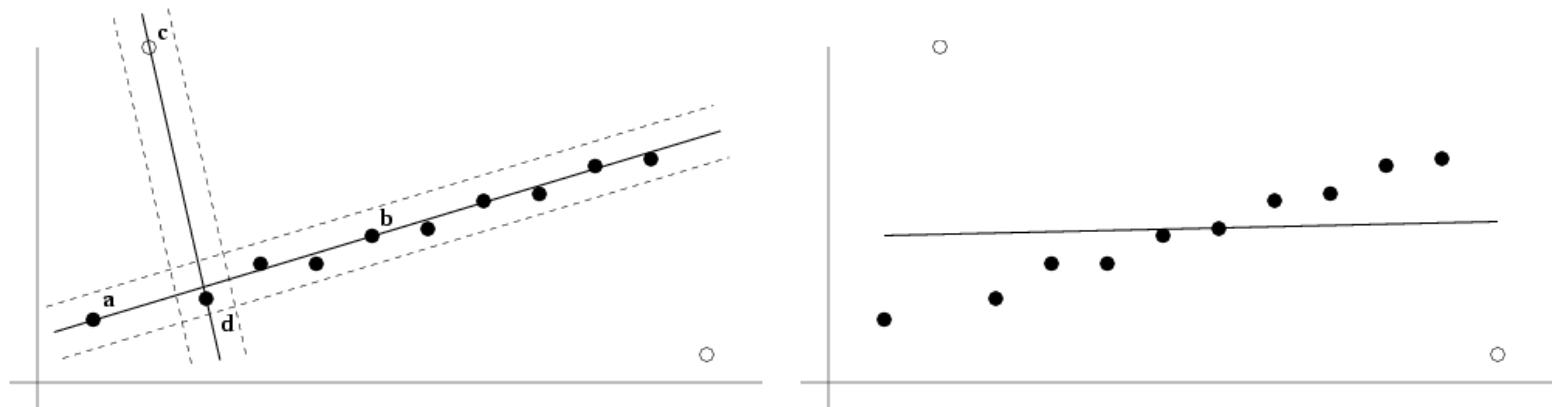
(for all points!)

# Recommendations:

1. Do not use unnormalized algorithms
2. Quick and easy to implement: 8-point normalized
3. Better: enforce rank-2 constraint during minimization
4. Best: Maximum Likelihood Estimation  
(minimal parameterization, sparse implementation)

# Robust Estimation

- What if set of matches contains gross outliers?



# RANSAC

## Objective

Robust fit of model to data set  $S$  which contains outliers

## Algorithm

- (i) Randomly select a sample of  $s$  data points from  $S$  and instantiate the model from this subset.
- (ii) Determine the set of data points  $S_i$  which are within a distance threshold  $t$  of the model. The set  $S_i$  is the consensus set of samples and defines the inliers of  $S$ .
- (iii) If the subset of  $S_i$  is greater than some threshold  $T$ , re-estimate the model using all the points in  $S_i$  and terminate
- (iv) If the size of  $S_i$  is less than  $T$ , select a new subset and repeat the above.
- (v) After  $N$  trials the largest consensus set  $S_i$  is selected, and the model is re-estimated using all the points in the subset  $S_i$

# How many samples?

- Choose  $t$  so probability for inlier is  $\alpha$  (e.g. 0.9)
  - Or empirically
- Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers. e.g.  $p = 0.99$

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

$s$	proportion of outliers $e$							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

# Acceptable consensus set?

- Typically, terminate when inlier ratio reaches expected ratio of inliers

$$T = (1 - e)n$$

# Adaptively determining the number of samples

$e$  is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield  $e=0.2$

- $N=\infty$ ,  $\text{sample\_count}=0$
- While  $N > \text{sample\_count}$  repeat
  - Choose a sample and count the number of inliers
  - Set  $e = 1 - (\text{number of inliers}) / (\text{total number of points})$
  - Recompute  $N$  from  $e$
  - Increment the  $\text{sample\_count}$  by 1
- Terminate

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

# RANSAC for F Estimation

Step 1. Extract features

Step 2. Compute a set of potential matches

Step 3. do

Step 3.1 select minimal sample (i.e. 7 matches)  
Step 3.2 compute solution(s) for F

Step 3.3 determine inliers (verify hypothesis)

until  $p(\#inliers, \#samples) > 95\% \text{ or } 99\%$

Step 4. Compute F based on all inliers

Step 5. Look for additional matches

Step 6. Refine F based on all correct matches

$$p = 1 - \left(1 - \left(\frac{\#inliers}{\#matches}\right)^7\right)^{\#samples}$$

#inliers	90%	80%	70%	60%	50%
#samples	5	13	35	106	382

# Finding more matches



restrict search range to neighborhood of epipolar line  
( $\pm 1.5$  pixels)

relax disparity restriction (along epipolar line)

# Degenerate Cases

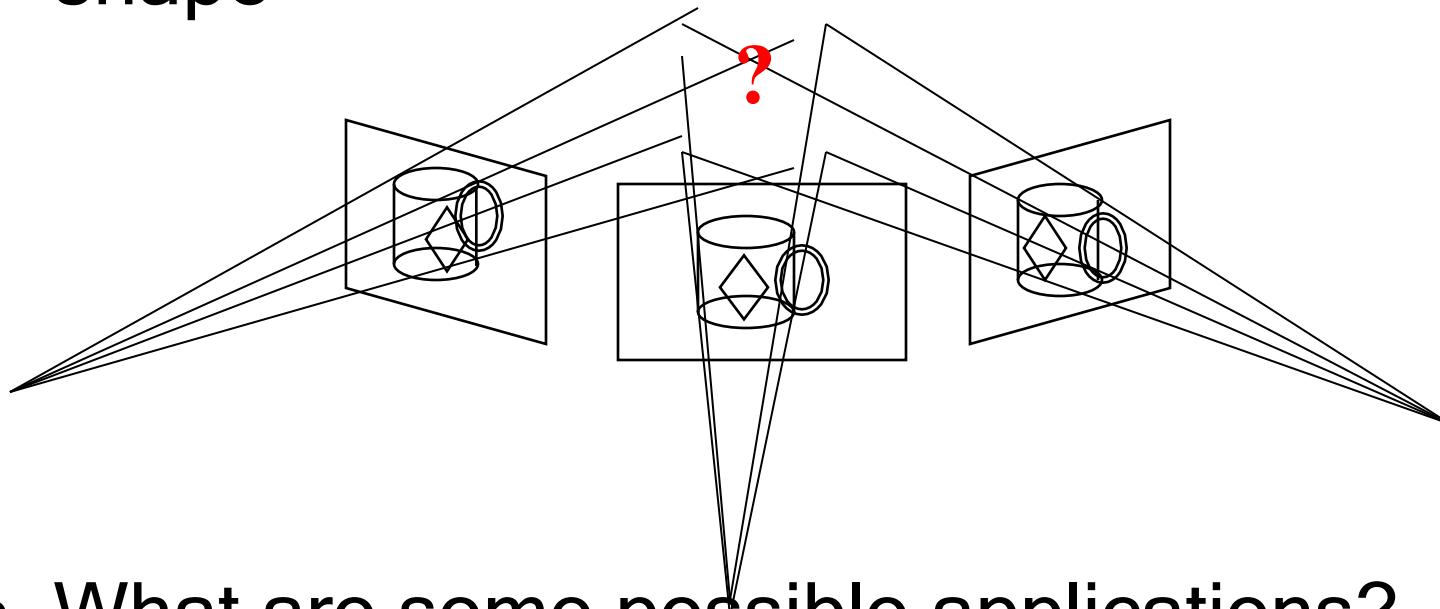
- Degenerate cases
  - Planar scene
  - Pure rotation
- No unique solution
  - Remaining DOF filled by noise
  - Use simpler model (e.g. homography)
- Model selection (Torr et al., ICCV'98, Kanatani, Akaike)
  - Compare H and F according to expected residual error (compensate for model complexity)

# Stereo Matching

Slides by Rick Szeliski, Pascal  
Fua and P. Mordohai

# Stereo Matching

- Given two or more images of the same scene or object, compute a representation of its shape



- What are some possible applications?

# Stereo Matching

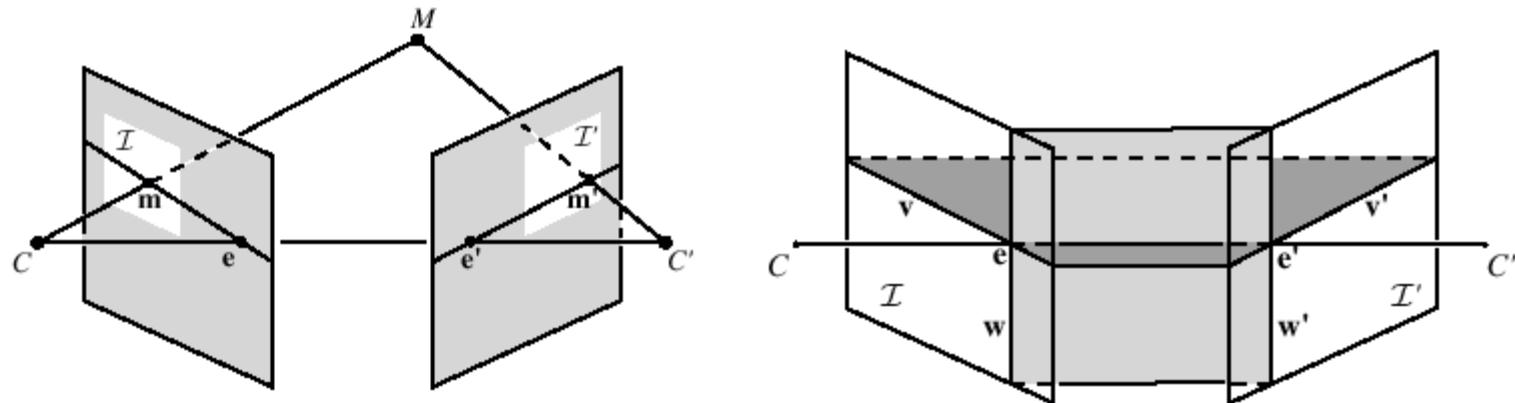
- Given two or more images of the same scene or object, compute a representation of its shape
- What are some possible representations?
  - depth maps
  - volumetric models
  - 3D surface models
  - planar (or offset) layers

# Stereo Matching

- What are some possible algorithms?
  - match “features” and interpolate
  - match edges and interpolate
  - match all pixels with windows (coarse-fine)
  - use optimization:
    - iterative updating
    - dynamic programming
    - energy minimization (regularization, stochastic)
    - graph algorithms

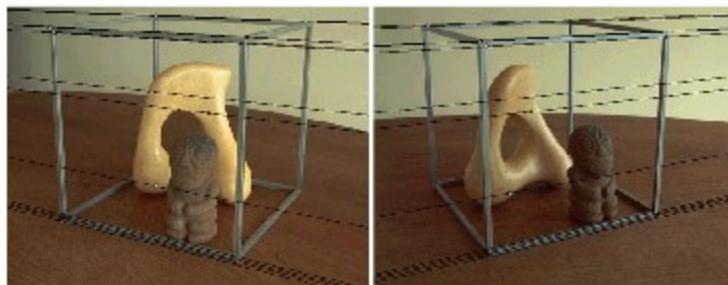
# Rectification

- Project each image onto same plane, which is parallel to the baseline
- Resample lines (and shear/stretch) to place lines in correspondence, and minimize distortion



- Take rectification for granted in this course

# Rectification



(a) Original image pair overlaid with several epipolar lines.



(b) Image pair transformed by the specialized projective mapping  $H_p$  and  $H'_p$ . Note that the epipolar lines are now parallel to each other in each image.

BAD!

# Rectification



(c) Image pair transformed by the similarity  $\mathbf{H}_r$  and  $\mathbf{H}'_r$ . Note that the image pair is now rectified (the epipolar lines are horizontally aligned).



(d) Final image rectification after shearing transform  $\mathbf{H}_s$  and  $\mathbf{H}'_s$ . Note that the image pair remains rectified, but the horizontal distortion is reduced.

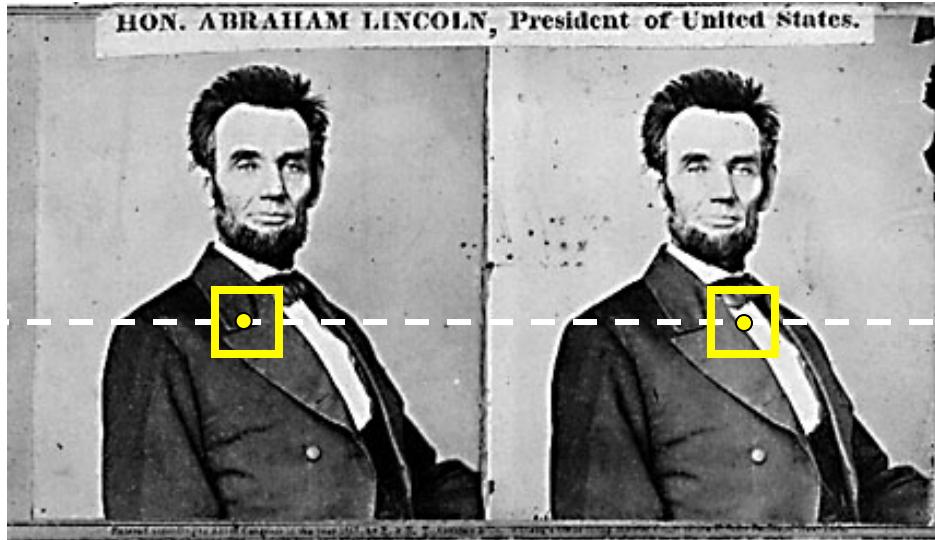
GOOD!

# Finding Correspondences

- Apply feature matching criterion at *all* pixels simultaneously
- Search only over epipolar lines (many fewer candidate positions)



# Basic Stereo Algorithm



For each epipolar line

For each pixel in the left image

- compare with every pixel on same epipolar line in right image
- pick pixel with minimum match cost

Improvement: match **windows**

# Disparity

- Disparity  $d$  is the difference between the  $x$  coordinates of corresponding pixels in the left and right image

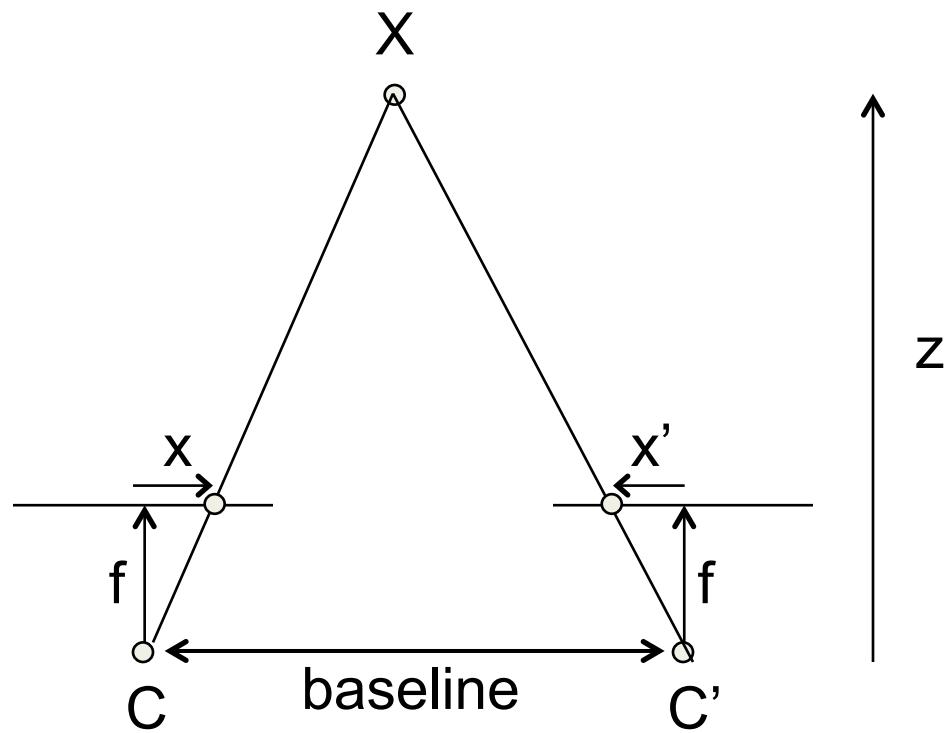
$$d = x_L - x_R$$

- Disparity is inversely proportional to depth

$$Z = bf/d$$

# Stereo Reconstruction

$$Z = bf/d$$



# Finding Correspondences

- How do we determine correspondences?
  - *block matching* or *SSD* (sum squared differences)

$$SSD = \sum_{[i,j] \in R} (f(i,j) - g(i,j))^2$$

- $d$  is the *disparity* (horizontal motion)



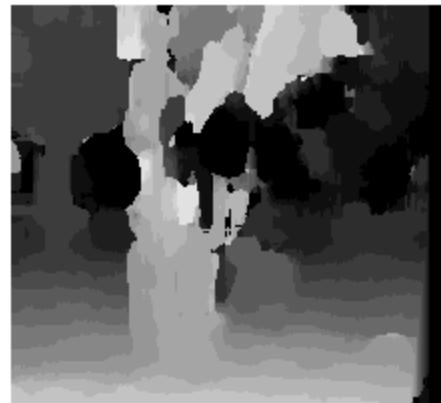
- How big should the neighborhood be?

# Neighborhood size

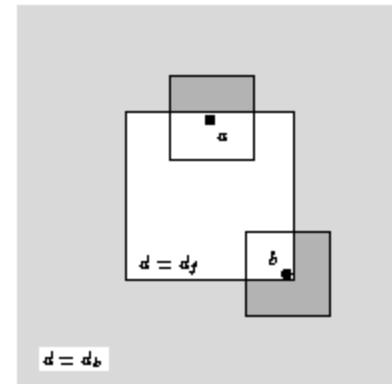
- Smaller neighborhood: more details
- Larger neighborhood: fewer isolated mistakes



$w = 3$



$w = 20$



# Challenges

- Ill-posed inverse problem
  - Recover 3-D structure from 2-D information
- Difficulties
  - Uniform regions
  - Half-occluded pixels
  - Repeated patterns



# Pixel Dissimilarity

- Sum of Squared Differences of intensities (SSD)

$$SSD = \sum_{[i,j] \in R} (f(i, j) - g(i, j))^2$$

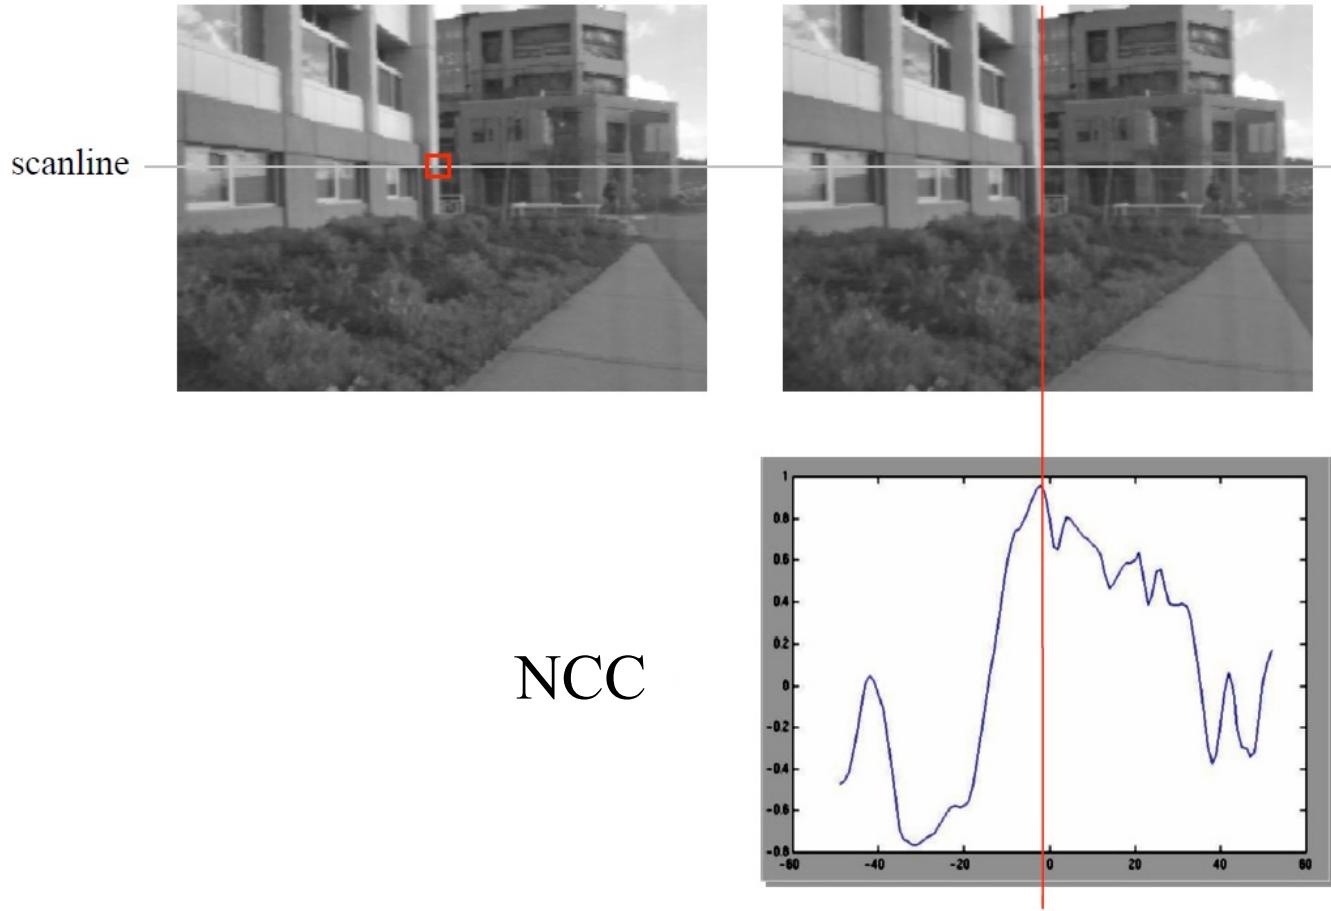
- Sum of Absolute Differences of intensities (SAD)

$$SAD = \sum_{[i,j] \in R} |f(i, j) - g(i, j)|$$

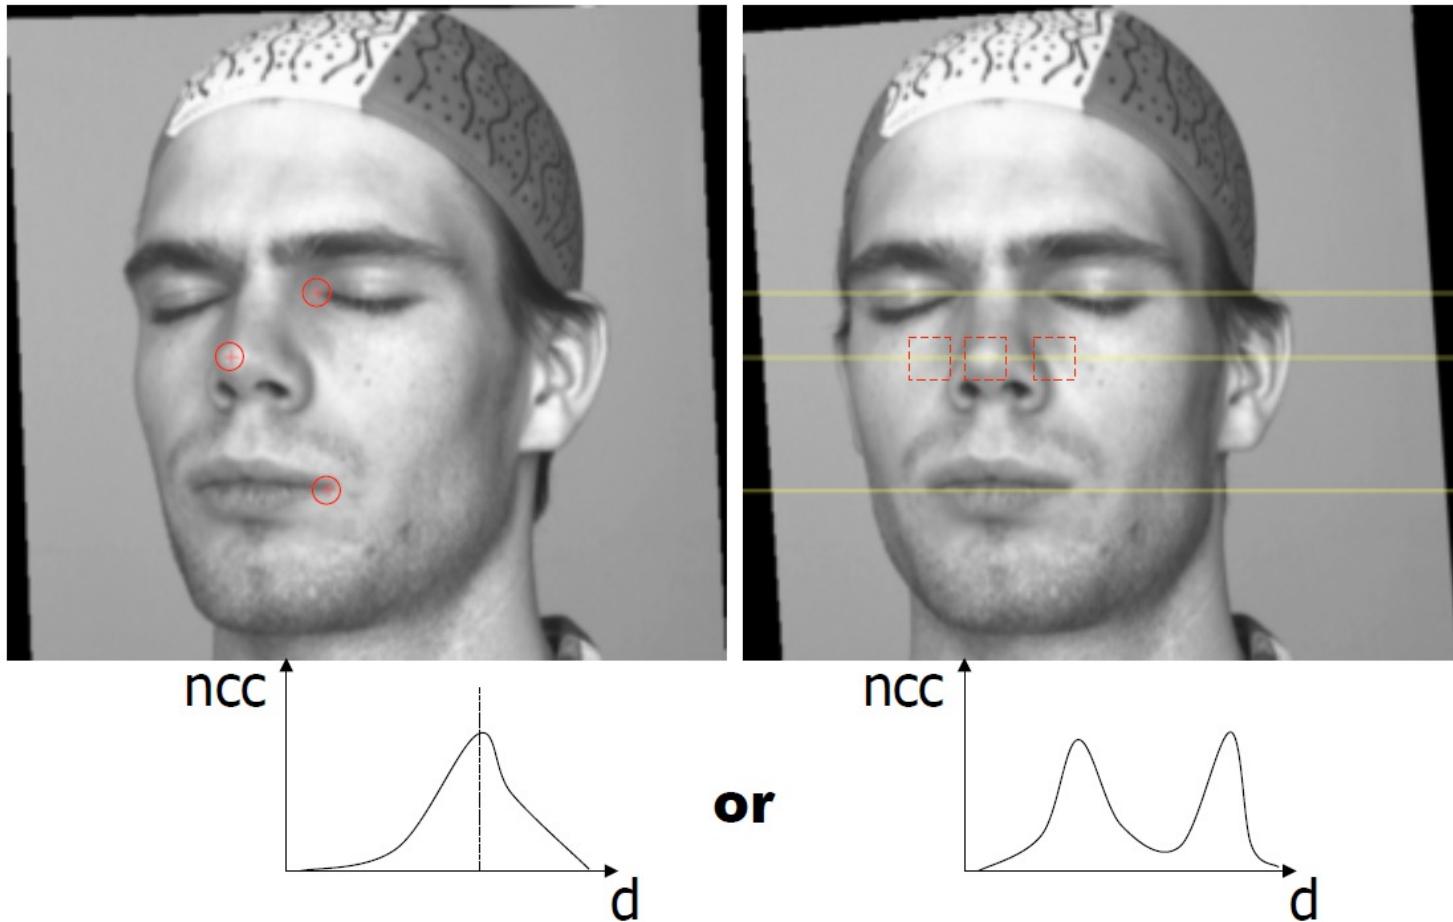
- Zero-mean Normalized Cross-correlation (NCC)

$$NCC(x, y, d) = \frac{\sum_{i \in W} (I_L(x_i, y_i) - \mu_L)(I_R(x_i - d, y_i) - \mu_R)}{\sigma_L \sigma_R}$$

# Cost/Score Curve

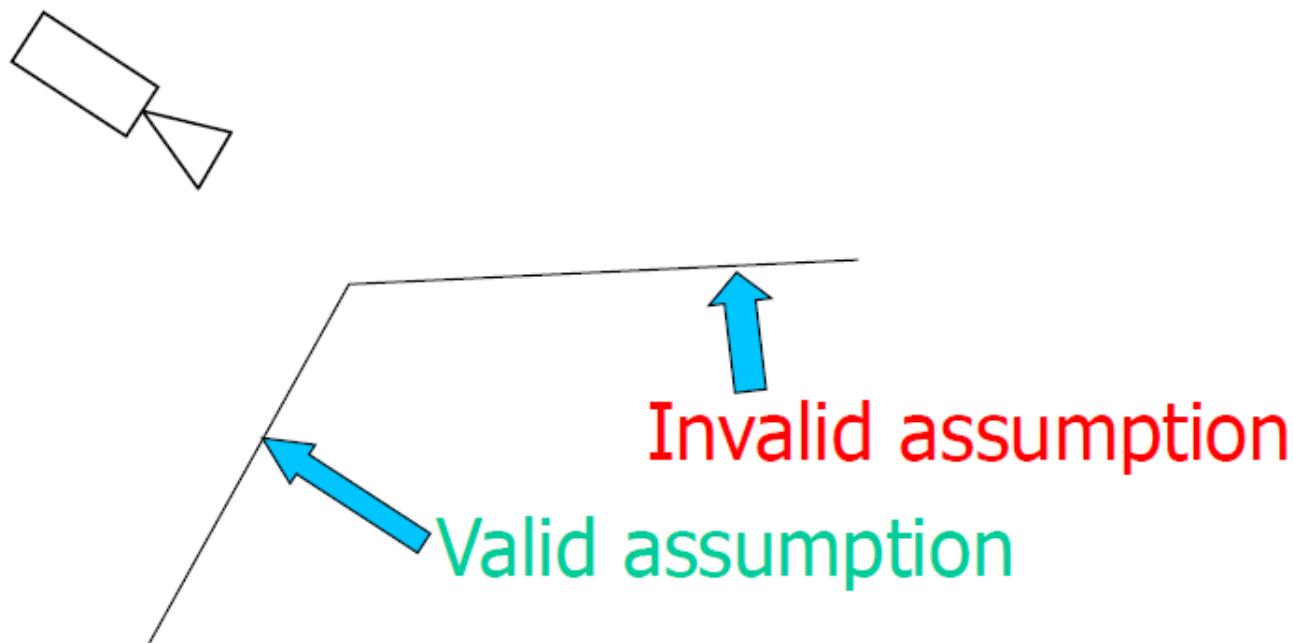


# Cost/Score Curve



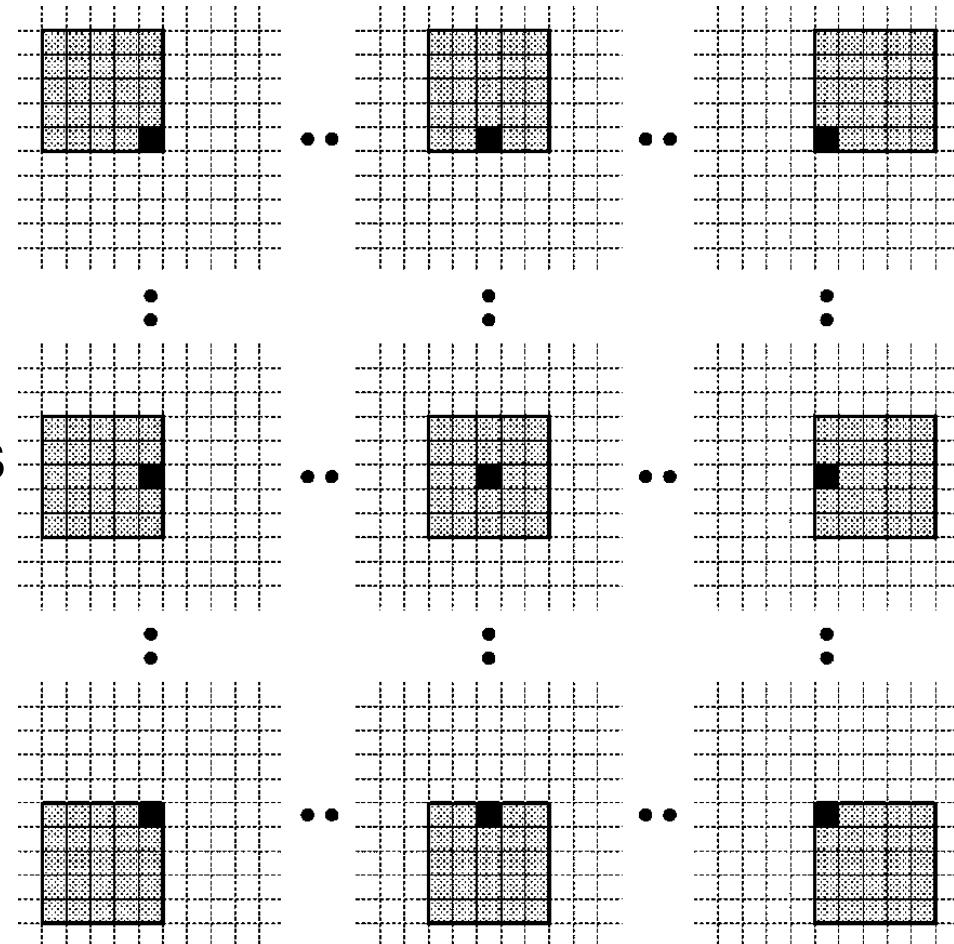
# Fronto-Parallel Assumption

- The disparity is assumed to be the same in the entire matching window
  - equivalent to assuming constant depth



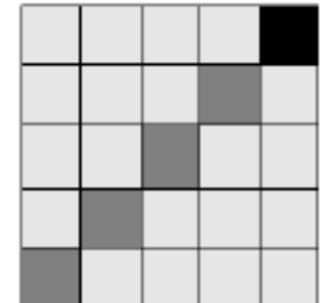
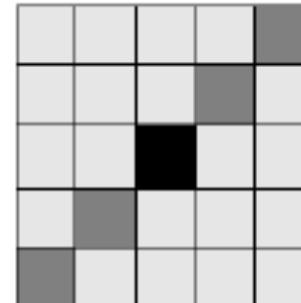
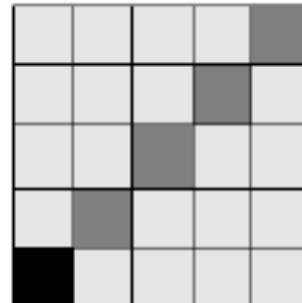
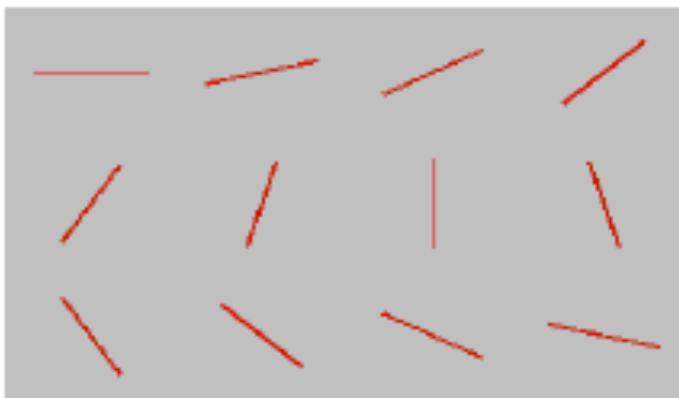
# Shiftable Windows

- Avoid having using matching windows that straddle two surfaces
  - Disparity will not be constant for all pixels
- Shift the window around the reference pixel
  - Keep the one with min cost (max NCC)



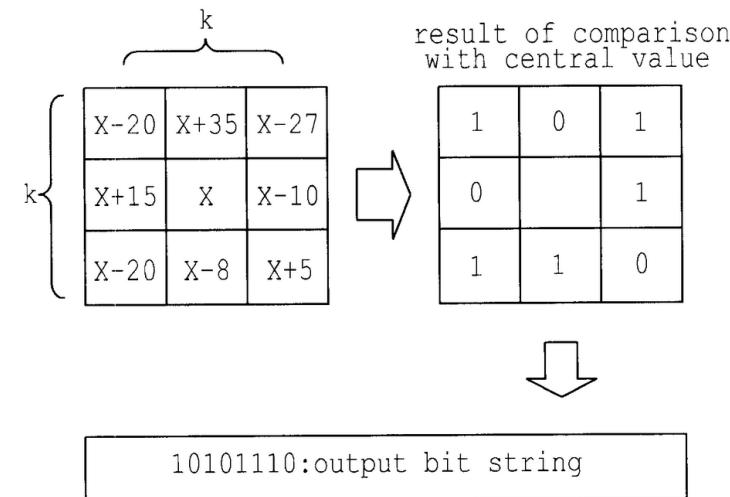
# Rod-shaped Filters

- Instead of square windows aggregate cost in rod-shaped shiftable windows
- Search for one that minimizes the cost (assume that it is an iso-disparity curve)



# Alternative Dissimilarity Measures

- Rank and Census transforms
- Rank transform:
  - Define window containing R pixels around each pixel
  - Count the number of pixels with lower intensities than center pixel in the window
  - Replace intensity with rank (0..R-1)
  - Compute SAD on rank-transformed images
- Census transform:
  - Use bit string, defined by neighbors, instead of scalar rank
- Robust against illumination changes



# Locally Adaptive Support

Apply weights to contributions of neighboring pixels according to **similarity** and **proximity**



(a) left support window  
(b) right support window  
(c) color difference  
between (a) and (b)

# Locally Adaptive Support

- Similarity in CIE Lab color space:

$$\Delta c_{pq} = \sqrt{(L_p - L_q)^2 + (a_p - a_q)^2 + (b_p - b_q)^2}$$

- Proximity: Euclidean distance

- Weights:  $w(p, q) = k \cdot \exp \left( -\left( \frac{\Delta c_{pq}}{\gamma_c} + \frac{\Delta g_{pq}}{\gamma_p} \right) \right)$

# Locally Adaptive Support: Results



(a) left image



(b) ground truth



(c) shiftable win. [7]



(d) compact win. [3]



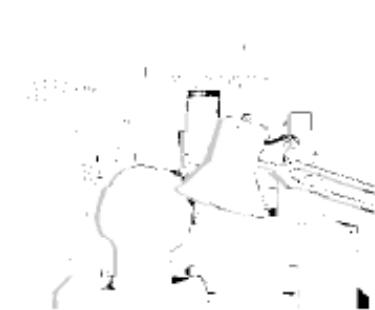
(e) variable win. [4]



(f) Bay. diff. [19]



(g) our result



(h) bad pixels (error > 1)